

Statistics for linguists

2023-12-06

Chi-squared test, loglinear models

From last class

- Linear models
- Errors and residuals
- Sum of squares
- ANOVA

For this class

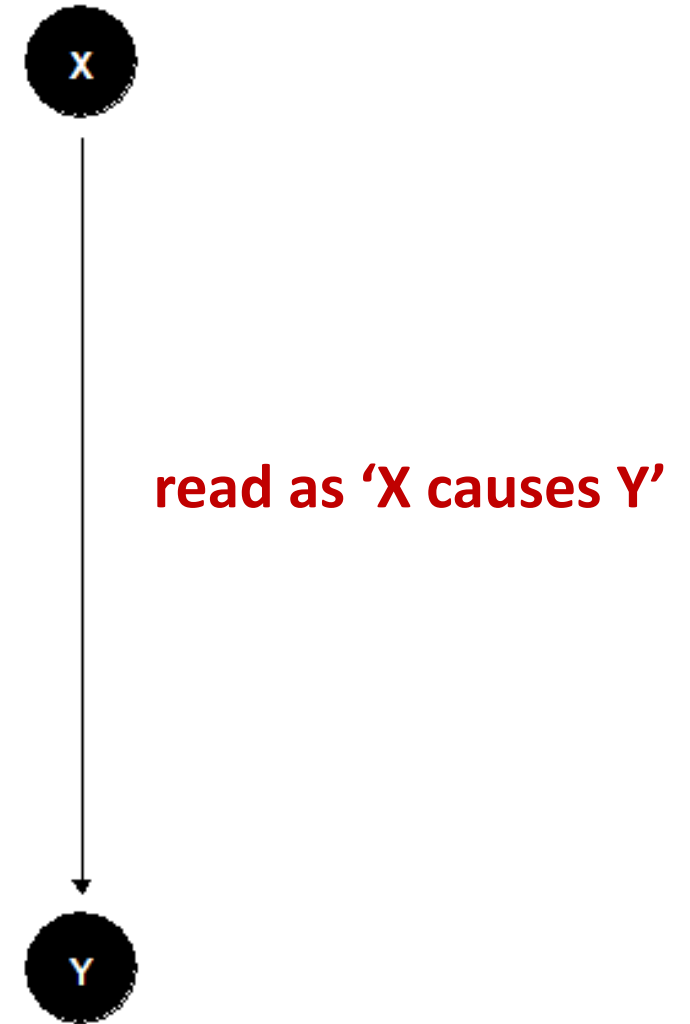
- Directed acyclic graphs (introduction)
- Chi-squared test
- Loglinear model / logistic regression
- Some data wrangling (splitstr etc.)

Packages for today

```
library(tidyverse)  
library(ggdag)  
library(V8)  
library(dagitty)  
library(glm2)  
library(Rling)  
library(rms)  
library(visreg)  
library(car)
```


Direct acyclic graph


- A directed acyclic graph represents our assumptions about causal direction.
- When you use a statistical model you imply assumptions about causation between your variables.
- So far, we have just been dealing with very simple causal relations (simple 'bivariate' models)




Directed acyclic graph (DAG)

- The causal relations can get very complicated though.
- DAGs are used for **causal inference** – if your statistical model is paired with DAG it means you can infer causation from your model, rather than just correlation





Causal inference & the Causal revolution

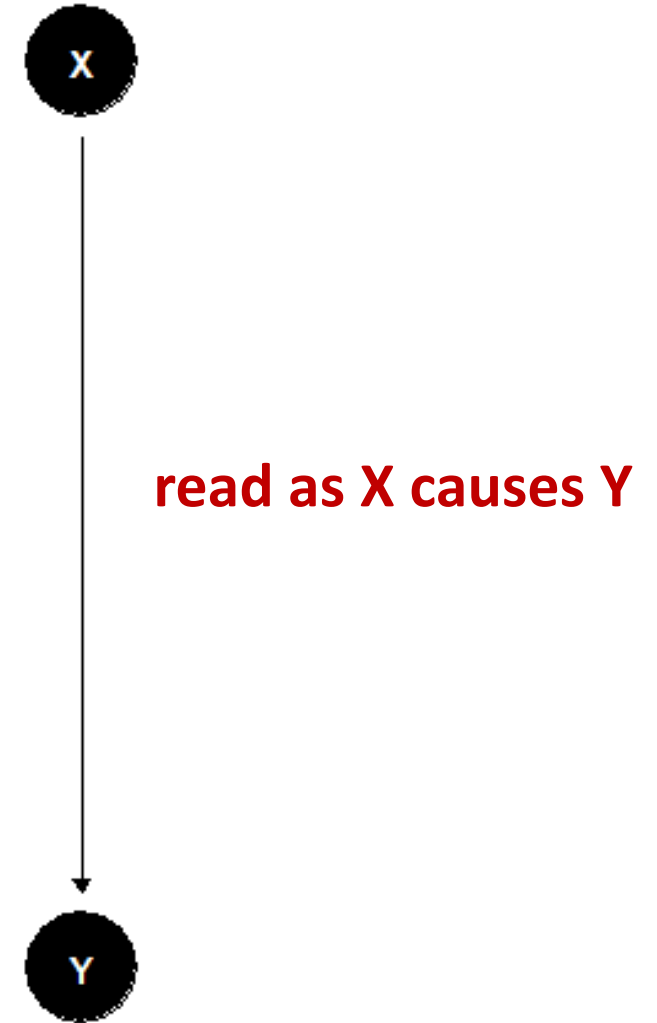


Pearl, Judea. 2009. *Causality: Models, Reasoning and Inference*. Cambridge University Press.

Pearl, Judea (with Dana Mackenzie). *The book of Why: The new science of cause and effect*. Penguin.

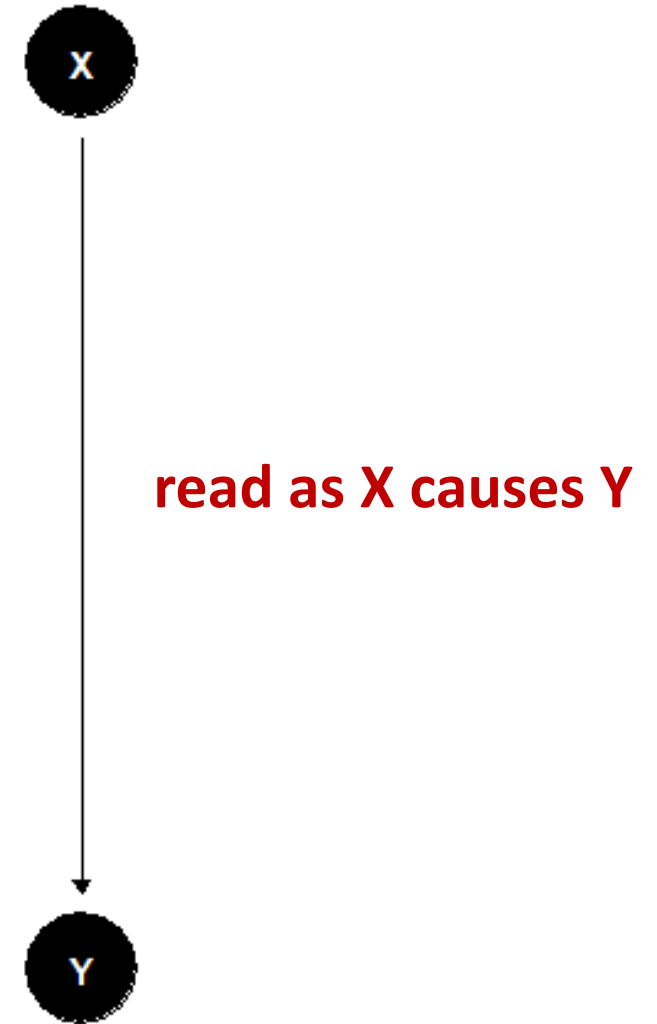
Direct acyclic graph

- Linear model / Regression
 - X = continuous
 - Y = continuous
- ANOVA
 - X is categorical
 - Y is continuous




Direct acyclic graph

- Linear model / Regression
 - X = continuous
 - Y = continuous
- ANOVA
 - X is categorical
 - Y is continuous
- Chi-squared
 - X is categorical
 - Y is continuous
- Loglinear model / Logistic regression
 - X is continuous
 - Y is categorical



Directed acyclic graph

- Naturally there can be more than one cause
- Next week we will start looking at *multivariate* models in more detail



Counts & contingency tables

- A lot of statistical information comes from counts (e.g. frequency of words in different texts)
- The data are usually presented in a contingency table.
- Data from Matthew Dryer (1992)
- OV = Object-Verb words order / VO = Verb-Object word order
- Postp = positions / Prep = prepositions
- H1: postpositions are associated with OV word order and prepositions are associated with VO order

Contingency table (word order associations)

	OV	VO
Postp	107	12
Prep	7	70

THE GREENBERGIAN WORD ORDER CORRELATIONS
MATTHEW S. DRYER
State University of New York at Buffalo

```
adpos <- matrix(c(107,12,7,70),ncol=2,byrow=TRUE)
rownames(adpos)<-c("PostP","Prep")
colnames(adpos)<-c("OV","VO")
adpos
```

Counts and probabilities

- These are **observed frequencies**
- We now need a model that predicts the **expected frequencies**
- Using these data, what is the probability of a random language from this sample having OV?

Counts and probability

```
wordorder <- cbind(c(107, 7), c(12, 70))
rownames(wordorder) <- c("Postp", "Prep")
colnames(wordorder) <- c("OV", "VO")
wordorder <- rbind(wordorder, c(114, 82))
wordorder <- cbind(wordorder, c(119, 77, 196))
rownames(wordorder) <- c("PostP", "Prep", "Column Total")
colnames(wordorder) <- c("OV", "VO", "Row total")
wordorder
```

Expected Frequency

- Raw total (Postp = 119, Prep = 77)
- Column total (OV = 114, VO = 82)
- Grand total = 196
- Expected = (Raw total * Column total) / Grand total

Crawley, Michael J. 2015. *Statistics: An Introduction using R*. Wiley.

Dryer, Matthew S. 1992. The Greenbergian Word Order Correlations. *Languages* 68:1, 81-138.

Expected frequency

- The expected frequency refers to what the values would be if VO/OV and Postp/Prep were independent.

	OV	VO
Postp	$(114 \cdot 119) / 196$	$(82 \cdot 119) / 196$
Prep	$(114 \cdot 77) / 196$	$(82 \cdot 77) / 196$

```
E <- cbind(c((114*119)/196, (114*77)/196),  
           c(82*119/196, (82*77)/196))  
rownames(E) <- c("Postp", "Prep")  
colnames(E) <- c("OV", "VO")  
E
```


Expected frequency

- What we've done is created a hypothetical “null distribution” against which we can measure how surprising our actual data are.

	OV	VO
Postp	69.21429	49.78571
Prep	44.78571	32.21429

Expected frequency vs. real frequencies

- What we've done is created a hypothetical "null distribution" against which we can measure how surprising our actual data are.

	OV	VO
Postp	69.21429	49.78571
Prep	44.78571	32.21429

	OV	VO
Postp	107	12
Prep	7	70

Chi-squared test

- The classical way of doing this is Karl Pearson's chi-squared test.

$$\chi^2 = \sum \frac{(\textit{Observed} - \textit{Expected})^2}{\textit{Expected}}$$



https://upload.wikimedia.org/wikipedia/commons/2/21/Karl_Pearson_2.jpg

```
E <- cbind(c((114*119)/196, (114*77)/196),
           c(82*119/196, (82*77)/196))
rownames(E) <- c("Postp", "Prep")
colnames(E) <- c("OV", "VO")
E
```

$$Expected = \frac{Rowtotal * Columntotal}{Grandtotal}$$

```
E.df <- melt(E)
colnames(E.df) <- c("Adposition", "Verb.Object",
                  "Expected.Frequency")
E.df$Observed.Frequency <- c(107, 7, 12, 70)
E.df
```

**Putting expected and observed
data in the same data set**

```
E.df$oe <- ((E.df$Observed.Frequency - E.df$Expected.Frequency)^2) / E.df$Expected.Frequency
sum(E.df$oe)
```


$$\chi^2 = \sum \frac{(Observed - Expected)^2}{Expected}$$

Chi-squared test and p values


- Is this number big?
- What is the critical value of the chi-squared test?
- To calculate this, we need the degrees of freedom and the cut off area you want.
- R = number of rows
- C = number of columns

Chi-squared test and p-value

- We have another hypothetical distribution based like the t distribution.



Logistic regression



Logistic regression

- The logistic regression is a method for modelling binary data.
- The basic ideas can be extended to non-binary data as long as they are organized into levels.
- It is typically used when the dependent variable is binary and there is an interest in knowing how a change in x effects the probability that something is y .

Logistic regression

- Psycholinguistic experiments where subjects have to give yes or no answers.
- Various uses in natural language processing
- Predict the risk of developing a specific disease
- Predict the probability that someone will vote for a particular political party.
- ...

Field experiment

- Let's look at data with a binary response variable.
- A **multiple forced choice** test was used with Chácobo speakers in 2022 using Praat
- Pitch and duration values of a form *janáquë* [hanáki] were manipulated systematically to create a matrix of duration and pitch values.
- Speakers Heard these and had to guess whether they were hearing *jánaquë* 's/he left it' or *janáquë* 's/he vomited'

Multiple forced choice experiment

- You'll have to wrangle the data a bit.

```
mcf <- read.csv("YourPathway/chacobo.mcf.df.csv")
response <- ifelse(mcf$response=="janáquë vomitó", 1, 0)
stimulus <- str_remove_all(mcf$stimulus, "janaquë_")
df <- as.data.frame(cbind(response, stimulus))
df[,c("Pitch_Hz", "Duration_ms")] <- str_split_fixed(df$stimulus, "_", 2)
df$Pitch_Hz <- as.numeric(df$Pitch_Hz)
df$response <- as.numeric(df$response)
head(df)
```

```
mcf <- read.csv("YourPathway/chacobo.mcf.df.csv")
response <- ifelse(mcf$response=="janáquë vomitó", 1, 0)

stimulus <- str_remove_all(mcf$stimulus, "janaquë_")
df <- as.data.frame(cbind(response, stimulus))
df[,c("Pitch_Hz", "Duration_ms")] <- str_split_fixed(df$stimulus, "_", 2)

df$Pitch_Hz<-as.numeric(df$Pitch_Hz)
df$response <- as.numeric(df$response)

head(df)
```

Change values so that they are 0 or 1

Extract the pitch and duration values from the stimuli

Make the values numeric

Linear model


- You could run a normal linear model on these data once the response variable is coded as 0 or 1.
- It is not very realistic or well-fit – but let's illustrate this anyways.

```
ggplot(df, aes(jitter(Pitch_Hz), jitter(response))) +  
  geom_point()+  
  stat_smooth(method="lm",  
              formula = y ~ x,  
              geom= "smooth")+  
  ylim(0,1)
```


Note that we use *jitter()* when lots of data points are overlapping so we can see the distribution.

Normal linear model

- You can see that as the pitch increases on the second vowel, the probability that *janáquë* 'vomit' = 1 is chosen over *jánaquë* 'leave' = 0 increases.
- But the linear model makes predictions way beyond the structure of the data.



```
plot(jitter(df$response)~jitter(df$Pitch_Hz), ylab="Response (0=accent first, 1=accent  
second)", xlab = "Frequency stimuli (Hz)", xlim=c(50,200), ylim=c(-2,2)) +  
abline(lm(response ~ Pitch_Hz, data = df))
```



Asymptotic function

- To model probability of one answer over another in relation to some data point x , we want a line that has **asymptotic** properties
- **Asymptotic function:** A function which gets closer and closer to a value or condition but never reaches it.
- Why? The difference between 125 and 150 Hz should be huge in terms of predicting the response but the difference between 150 and 175 Hz should be smaller.

Logistic regression

- A logistic regression or logit model can be represented with the following equations

$$\text{logit}(y) = \alpha + \beta_1 x_1 + \beta_2 x_2 \dots + e$$

$$\text{logit}(p) = \log \frac{p}{1-p}$$

$$\text{Prob}\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

Logistic regression

- A logistic regression or logit model can be represented with the following equations

$$\textit{logit}(y) = \alpha + \beta_1 x_1 + \beta_2 x_2 \dots + e$$

$$\textit{logit}(p) = \log \frac{p}{1-p}$$

$$\textit{Prob}\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

The intercept plus the coefficient β times x (plus error etc.) is equal to y output from a logit function

Logistic regression

- A logistic regression or logit model can be represented with the following equations

$$\text{logit}(y) = \alpha + \beta_1 x_1 + \beta_2 x_2 \dots + e$$

$$\text{logit}(p) = \log \frac{p}{1-p}$$

$$\text{Prob}\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

The logit function is the log of a probability divided by 1 minus that probability – the “log odds”

Logistic regression

- A logistic regression or logit model can be represented with the following equations

$$\text{logit}(y) = \alpha + \beta_1 x_1 + \beta_2 x_2 \dots + e$$

$$\text{logit}(p) = \log \frac{p}{1-p}$$

$$\text{Prob}\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

So to get the probability that y is 1 given x, you divide the loglinear variables by an equation that reverses the logit function.

Logistic regression

- A logistic regression or logit model can be represented with the following equations

$$\text{logit}(y) = \alpha + \beta_1 x_1 + \beta_2 x_2 \dots + e$$

Just algebra

$$\text{logit}(p) = \log \frac{p}{1-p}$$


$$\text{Prob}\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

So to get the probability that y is 1 given x, you divide the loglinear variables by an equation that reverses the logit function.

Bernoulli distribution

- A Bernoulli distribution is a **discrete distribution** with **two possible outcomes**.
- We want to account for the distribution of y .
- You can model it with `rbinom()`

$$P(y) = \begin{cases} 1 - p, & \text{for } y = 0 \\ p, & \text{for } y = 1 \end{cases}$$



By Niklaus Bernoulli (1662-1716) - [2] [3], Public Domain,
<https://commons.wikimedia.org/w/index.php?curid=266673>

Bernoulli distribution

- A Bernoulli distribution is a discrete distribution with two possible outcomes.
- We want to account for the distribution of y .
- You can model it with `rbinom()`

$$P(y) = \begin{cases} 1 - p, & \text{for } y = 0 \\ p, & \text{for } y = 1 \end{cases}$$

‘The probability of y being 1 is p and the probability of y being 0 is $p - 1$ ’

Bernoulli distribution

- A logistic regression or logit model can be represented with the following equations

$$P(y) = \begin{cases} 1 - p, & \text{for } y = 0 \\ p, & \text{for } y = 1 \end{cases}$$

```
bernouilli_data <- rbinom(20, 1, 0.5)
print(bernouilli_data)
## [1] 1 1 1 0 0 1 0 0 1 0 0 1 0 0 0 1 0 0 0 1
```


Logistic function

$$\text{logit}(p) = \log \frac{p}{1-p}$$

- The important point about log odds ratios is that they take any numbers and transform them to a number from 0 to 1 along a sigmoid shape.
- Why is this good?
- Because we are interested in modelling a binary outcome, 0 or 1, and we want to have a function that translates the effect of predictors into that scale for y .


Logistic function

- Any value of x will be transformed into a number that varies from 0 to 1 with ceiling effects.

$$Prob\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

Logistic regression

- y is bounded to 0 or 1
- The relationship between x and y has a ceiling effect (like logarithms)
- Let's run a simulation model to get the feel for it.



Functions

- There are two functions in R
- The logarithmic function and the exponential function – they are basically mirror-images of each other.

```
t <- log(10)
t
## [1] 2.302585
exp(t)
## [1] 10
```

- The *odds* function is pretty straightforward


$$\frac{1}{1 - p}$$

Odds

- If you plot probabilities versus odds, you can see that odds has an asymptotic property in relation to probability


```
par(mfrow=c(1,1))  
x <- seq(0,1,by=0.1)  
plot(x, 1/(1-x))
```

$$\frac{1}{1-p}$$




Plotting odds

```
plot(x, 1/(1-x))
```




```
plot(1/(1-x), x)
```




Logs

- Logs also have this asymptotic property on both sides.

```
x <- seq(0,1,by=0.1)
plot(x,log(x))
```



```
x <- seq(-2,0, by =.1)
plot(x, exp(x))
```





- But not exactly how we want because we want a ceiling effect approaching 0 and approaching 1

Log-odds

```
par(mfrow=c(1,1))  
x <- seq(0,1,by=.02)  
plot(log(x/(1-x)),x)
```

```
par(mfrow=c(1,1))  
x <- seq(-4,4,by=.1)  
plot(x, 1/(1+exp(-x)))
```



Slope of an S-curve





- Does a slope make sense?
- The slope changes as x changes – so there's not a single slope.
- There are concepts we can make use of (e.g. the slope at the steepest point of the curve)
- It is hard to interpret coefficients in logistic regressions

Slope

- Adding a slope coefficient

```
par(mfrow=c(2,2))
x <- seq(-4,4,by=.1)
b <- 2
plot(x, 1/(1+exp(-x*b)), main = "b = 2")
b <- 8
plot(x, 1/(1+exp(-x*b)), main = "b = 8")
b <- 0.5
plot(x, 1/(1+exp(-x*b)), main = "b = 0.5")
b <- -1
plot(x, 1/(1+exp(-x*b)), main = "b = -1")
```

Slope of a logistic function



- Intuitively, the slope corresponds to how steep the curve is around its center point
- Large slope = small change in x is equal to a large change in the probability that $y = 1$ around the middle of the S-curve (where the slope is changing the fastest)

Simulation exercise

- Simulate whether someone will go to class other not

```
set.seed(1)
coffee <- rnorm(100, 15, 5)
happiness <- rnorm(100, 10, 2)
a <- -5
b1 <- -1
b2 <- 1
```

```
xb <- a + b1 * happiness + b2 * coffee + rnorm(100, 0, 0.1)
p <- 1 / (1 + exp(-xb))
```

Simulation exercise

- Simulate whether someone will go to class other not

```
set.seed(1)
coffee <- rnorm(100, 15, 5)
happiness <- rnorm(100, 10, 2)
a <- -5
b1 <- -1
b2 <- 1
```

```
xb <- a + b1 * happiness + b2 * coffee + rnorm(100, 0, 0.1)
p <- 1 / (1 + exp(-xb))
```

$$Prob\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

We have to subject our y to the inverse of the logit function to get the actual probabilities


Simulation exercise

```
summary(p)
```

```
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## 0.0000497 0.0780173 0.7105522 0.5493012 0.9611723 0.9999948
```

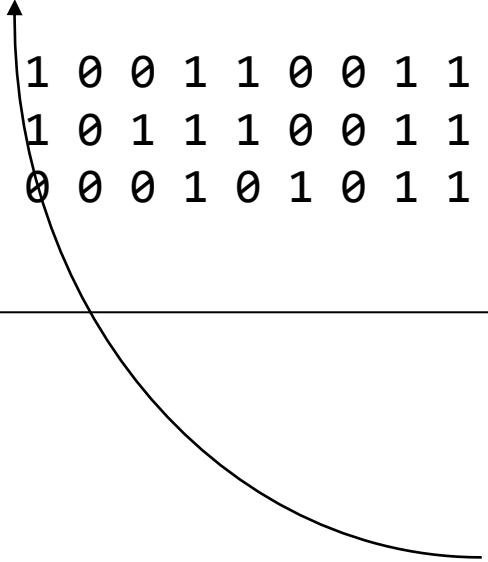
```
par(mfrow=c(1,1))
```

```
plot(p)
```



Simulation exercise

```
gotoclass <- rbinom(n=100, size=1, prob=p)
gotoclass
##      [1] 0 0 0 1 0 0 1 1 0 0 1 1 0 0 1 1 1 1 1 1 1 1 0 1 0 0 0 0 1 1 1 1 1 0 1 0
##     [38] 1 1 1 1 0 1 1 1 0 0 1 1 1 1 0 1 0 1 1 0 0 1 0 1 0 1 0 0 0 0 1 1 1 0 0 1 0
##     [75] 0 1 0 0 0 0 1 0 1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 0 0 0 0
```



$$P(y) = \begin{cases} 1 - p, & \text{for } y = 0 \\ p, & \text{for } y = 1 \end{cases}$$

We are creating a Bernoulli distribution from our probabilities which are produced from the stochastic model we simulated earlier

Generalized linear models

- Logistic regressions are a type of **generalized linear models**
- Unlike normal linear models, glms do not use ordinary least squared (OLS) but rather use **Maximum Likelihood** to estimate parameters.
 - (We will discuss Maximum Likelihood later in in class)
- Use the package glm2()

```
library(glm2)
model.logit <- glm(gotoclass~happiness+coffee, family="binomial")
```



```
summary(model.logit)
```

```
##
```

```
## Call:
```

```
## glm(formula = gotoclass ~ happiness + coffee, family = "binomial")
```

```
##
```

```
## Coefficients:
```

```
##           Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept)  -4.1157      2.8897  -1.424  0.15438
```

```
## happiness    -1.6263      0.5221  -3.115  0.00184 **
```

```
## coffee        1.3600      0.3348   4.062 4.87e-05 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
## Null deviance: 137.186 on 99 degrees of freedom
```

```
## Residual deviance: 36.172 on 97 degrees of freedom
```

```
## AIC: 42.172
```

```
##
```

```
## Number of Fisher Scoring iterations: 8
```

Interpreting logistic regression coefficients

- It is hard to interpret logistic regression coefficients because the relationship is non-linear
- The intercept is interpreted assuming 0 for other predictors
 - But sometimes 0 is not interesting
 - Alternatively we can interpret the intercept at the center point
- Rather than consider a discrete change in x we can compute an approximation of the derivative of the logistic curve at the central value (where the relationship is steepest)
 - You get this by dividing the coefficient by 4

Interpreting logistic regression coefficients

- **Divide by 4 rule (this will give you the coefficient)**
- Dividing by 4 gives you the **maximum difference** in y corresponding to a **unit of difference** in x

```
-1.6263/4  
## [1] -0.406575
```

A decrease in 0.4 of happiness corresponds to the steepest rise in the probability of going to class where coffee is at its mean value

Interpreting logistic regression

- The intercept is what the logit probability of y is when all the predictors are 0.

```
invlogit <- function (x) {1/(1+exp(-x))}
invlogit(-4.1157)
## [1] 0.01605263
```

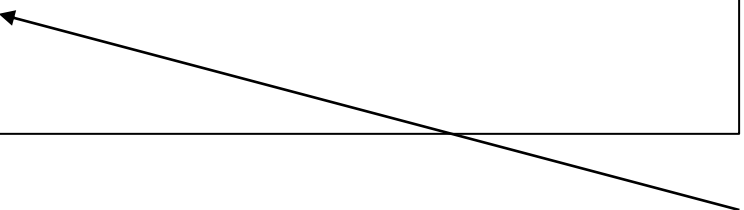
This is a function I wrote in R (actually in Gelman & Hill 2007), to calculate the probability from the logit probability

$$Prob\{y = 1|x\} = \frac{1}{1 + \exp(-x\beta)}$$

Interpreting logistic regression

- The intercept is what the logit probability of y is when all the predictors are 0.

```
invlogit <- function (x) {1/(1+exp(-x))}  
invlogit(-4.1157)  
## [1] 0.01605263
```



**If someone has 0 happiness score
and they have had no coffee, there is
a 1.6% chance they will go to class**

**The probability of y (going to class)
when happiness is 0 and coffee
consumption is 0.**

Interpreting logistic regression

- The intercept is what the logit probability of y is when all the predictors are 0.

```
invlogit <- function (x) {1/(1+exp(-x))}  
invlogit(-4.1157 + -1.6263*mean(happiness) + 1.3600*mean(coffee))  
## [1] 0.707797
```

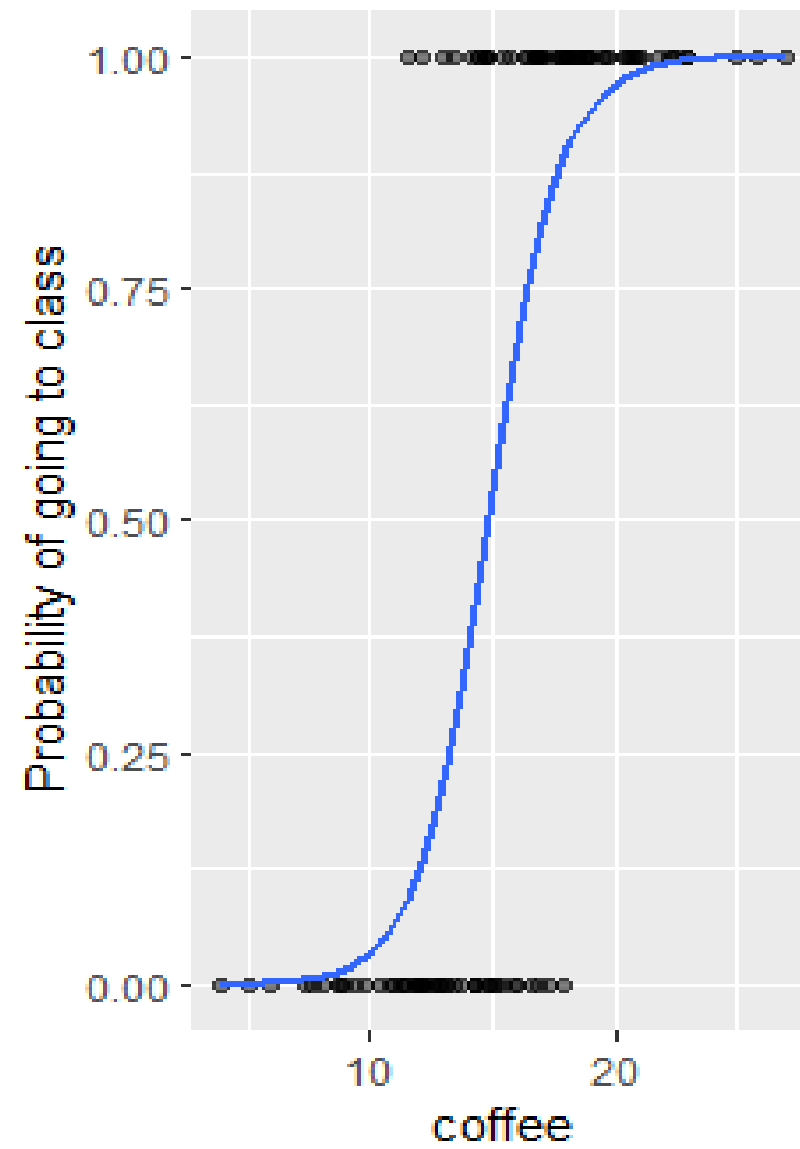

If someone has a mean happiness score and they have had the average amount of coffee, there is a 70.1% chance they will go to class

The probability of y (going to class) when happiness is at its mean value (-16.41) and coffee consumption is at its mean value (21.1)

Visualizing logistic regression in ggplot

```
par(mfrow=c(1,2))
data <- data.frame(gotoclass, happiness, coffee)
plotline<-ggplot(data, aes(x=coffee, y = gotoclass))+geom_point()+
  geom_abline(intercept = 0.11105, slope = 0.076446, color="red", size=1)+
  ylab("Probability of going to class")
```

```
plotS<-ggplot(data, aes(x=coffee, y= gotoclass))+
  geom_point(alpha=.5)+
  stat_smooth(method="glm", se=FALSE, method.args = list(family=binomial))+
  ylab("Probability of going to class")
plotS
```



Logistic regression on Chácobo MFC

- Try running a logistic regression model on the Chacobo Forced Experiment data, with the response as the dependent variable and the simulated Pitch (Hz) as the independent variable.

```
logit_model_01 <- glm(response~Pitch_Hz, data=df, family="binomial")  
summary(logit_model_01)
```

```
## Call:
## glm(formula = response ~ Pitch_Hz, family = "binomial", data = df)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -9.109549   0.626233  -14.55  <2e-16 ***
## Pitch_Hz     0.075658   0.005149   14.70  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2242.6  on 1619  degrees of freedom
## Residual deviance: 1980.1  on 1618  degrees of freedom
## AIC: 1984.1
##
## Number of Fisher Scoring iterations: 4
```

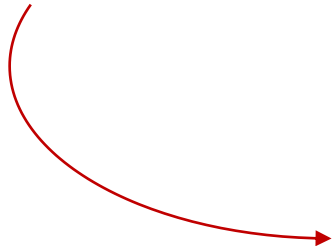
Interpreting the results

- Apply the Divide by 4 rule to the slope coefficient
- Interpret the intercept at the mean value for the Pitch (Hz)
- Plot the logistic regression using `ggplot()`

Interpreting the slope

0.075658/4

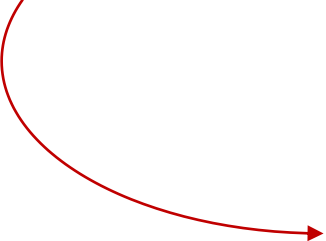
[1] 0.0189145



An increase in 0.02 Hz of pitch corresponds to the steepest rise in the probability of a Chácobo speaker choosing *janáquë* 's/he vomited' over *jánaquë* 's/he left' in the MCF

Interpreting the intercept


```
invlogit(-9.109549+ 0.075658*mean(df$Pitch_Hz))  
## [1] 0.5287489
```



The probability of *y* (a Chácobo speaker choosing *janáquë* 's/he vomited') when pitch (Hz) is at its mean for the experiment (121.93 Hz) is 52% (when the first syllable is fixed at 80 Hz)

Plotting the loglinear relation

```
plotS<-ggplot(dat=df, aes(x=Pitch_Hz, y= response))+  
  geom_point(alpha=.5)+  
  stat_smooth(method="glm", se=FALSE, method.args = list(family=binomial))+  
  ylab("Probability of choosing janáquë")  
plotS
```



Dutch causative constructions

- Predict which causative occurs in Dutch

```
data(doenLaten)
```

```
d <- doenLaten
```

```
head(d)
```

##	Aux	Country	Causation	EPTrans	EPTrans1
## 1	laten	NL	Inducive	Intr	Intr
## 2	laten	NL	Physical	Intr	Intr
## 3	laten	NL	Inducive	Tr	Tr
## 4	doen	BE	Affective	Intr	Intr
## 5	laten	NL	Inducive	Tr	Tr
## 6	laten	NL	Volitional	Intr	Intr

Dutch causative constructions

- Predict which causative occurs in Dutch

```
data(doenLaten)
d <- doenLaten
head(d)
```

##	Aux	Country	Causation	EPTrans	EPTrans1
## 1	laten	NI	Inducive	Intr	Intr

## 2	laten				
## 3	laten				
## 4	doen				
## 5	laten				
## 6	laten				

(1) *Hij deed me denken aan mijn vader*
He did me think at my father
'He reminded me of my father.'

(2) *Ik liet hem mijn huis schilderen*
I let him my house Paint
'I had him paint my house.'

***Doen* relates to direct causation**

***Laten* relates to indirect causation**